

When Algorithms Remember What We Forget

Computational Thinking as Infrastructure for Sustainable Community Service

Mei Hsien Hsu

Las Positas College Honors Program

CS 4: Introduction to Artificial Intelligence

Professor An Lam

Fall 2025

Abstract

This paper explores the application of seven fundamental artificial intelligence algorithms to volunteer coordination in community service organizations, using Tzu Chi San Francisco's Hunters Point operations as a case study. Drawing on sixteen years of personal volunteer experience and one semester of formal AI education, the author demonstrates how A* Search, Markov Decision Processes, Bayesian Networks, First-Order Logic, Propositional Logic, Adversarial Search, and Knowledge-Based Agents can serve as coordination infrastructure that respects volunteers' limited time while maximizing community impact.

The project culminates in "Journey of Kindness," an educational game that teaches these algorithms through humanitarian scenarios while incorporating the Ataraxy Portico (靜心之門)—a contemplative pause system featuring Jing Si Aphorisms that embodies the principle 「動中有靜， 靜中有動」 (within motion, find stillness; within stillness, find motion). Pilot testing with 15 participants demonstrated promising results: median ELO improvement of +180 points (IQR: 95-240) and 67% recruitment conversion rate (95% CI: 38%-88%), though these findings should be interpreted as preliminary given the small sample size.

Beyond technical implementation, this paper reflects on the ethical tensions that emerge when algorithmic optimization encounters human complexity: when should we trust the algorithm, and when must we override it for values that resist formalization? The final game level introduces a novel approach to AI ethics education, drawing on Mo Gawdat's philosophy that AI learns from what we teach it, using multi-AI dialogue to demonstrate how different value systems approach moral dilemmas.

The synthesis this paper proposes is that computational thinking does not replace compassion but provides the structure compassion needs to become sustainable service.

Keywords: *artificial intelligence, volunteer coordination, A* search, Markov decision processes, Bayesian networks, first-order logic, knowledge-based agents, adversarial search, AI ethics, humanitarian technology, gamification, Tzu Chi Foundation*



Figure 3. Journey of Kindness main menu showing seven AI algorithm levels through humanitarian scenarios.

Introduction: A Handful of Raw Rice

John Muir Elementary School, San Francisco, 2006. Every Friday morning, the San Francisco Food Bank delivered pallets of food to the school auditorium. Tzu Chi volunteers—members of an international Buddhist humanitarian organization founded in Taiwan—unpacked and sorted the provisions, assembling family-sized packages by 9 AM. Each package contained canned vegetables, daily necessities, cooking seasonings, and a bag of uncooked rice. Children picked them up before class; parents collected the groceries after school, taking home provisions that might constitute the week's primary food source.

One Friday morning, a volunteer named Roxanne witnessed something that would redirect her entire life trajectory. A young Black girl, perhaps seven years old, tore open her family's food bag on the spot. She reached inside, grabbed a fistful of raw, uncooked rice grains, and began chewing them.

Roxanne approached her gently. "Sweetheart, how long has it been since you've eaten?"

"Two days," the child replied. "Maybe three."

Roxanne—at that time a senior executive at Genentech with dual master's degrees in Computer Science and Finance, plus additional training in Psychiatry—made a decision that morning. She would not merely donate money or volunteer occasionally. She would restructure her entire professional life around systematic community service. Over the following nineteen years, she became a central figure in Tzu Chi San Francisco's operations, eventually becoming my mentor when I joined the organization in 2009.

This paper emerges from that lineage. In 2015, during my own health crisis—a kidney transplant at UC Davis where my daughter donated her kidney to me—I experienced the profound depth of Tzu Chi's coordinated compassion. After surgery, my husband accompanied my daughter and me, and Sacramento Tzu Chi volunteers Chou and Joe, an elderly couple whom I had never met before, drove us directly from UC Davis Medical Center back to our home on Hammonton Smartsville Road in Marysville. Though we had only connected through our shared identity as Tzu Chi dharma family (法親), they cared for

us with the warmth of lifelong friends. I experienced firsthand what it means to be on the receiving end of coordinated compassion.

Now, as a 58-year-old computer science student at Las Positas College, I am attempting to give back by applying what I've learned in Professor An Lam's AI course to the operational challenges I've observed over sixteen years of service.

The central question this paper addresses: Can classical AI algorithms—developed primarily for robotics, game-playing, and commercial optimization—serve humanitarian purposes? Can they help answer the question that time-constrained volunteers increasingly ask: "If I give you three hours of my limited monthly availability, will you coordinate those hours well enough to create genuine impact?"

My answer, after implementing these systems for Tzu Chi's operations: yes. But the implementation revealed something textbooks rarely discuss—the moments when algorithmic recommendations must be overridden because human situations resist complete formalization.

Background: Hunters Point and the Problem Space

The Community Context

Hunters Point sits at the southeastern edge of San Francisco, historically isolated from the city's wealth and services. Once home to a thriving Chinese shrimp-fishing community in the 1860s-1930s, the area was transformed during World War II into the Hunters Point Naval Shipyard. After the war, African American workers who had migrated for wartime employment remained, forming a community that would face decades of environmental contamination, economic disinvestment, and social marginalization.

Today, unemployment in Hunters Point exceeds 19%—far above San Francisco's average. The area contains some of the city's last remaining public housing, scattered across complexes that concentrate poverty and limit residents' access to the opportunities that define the city just miles away. Environmental justice concerns persist: the naval shipyard cleanup continues decades after closure, and residents report elevated rates of cancer and respiratory illness.

Tzu Chi San Francisco has maintained continuous presence in Hunters Point since 2009, operating through what we call the "Three Happinesses" (三福) framework:

Table 1

Tzu Chi San Francisco "Three Happinesses" (三福) Framework

Program	Core Activities	Frequency
Happy Campus (幸福校園)	Attendance incentives, supplies, summer programs, scholarships	Ongoing + seasonal
Happy Family (幸福家庭)	Food bank, home furnishings, employment support, winter care	Weekly + monthly
Happy Community (幸福社區)	Mobile dental, vision screening, health education, 3R workshops	Monthly

Happy Campus (幸福校園) partners with San Francisco Unified School District to support thousands of K-12 students across 33 schools. While in-person tutoring paused post-COVID, core activities remain active: attendance incentive programs reward perfect attendance with certificates and gifts; annual back-to-school events distribute approximately

1,600 uniforms along with backpacks, supplies, and health screenings; summer programs offer life exploration workshops emphasizing humanistic education; and scholarships support low-income high-achieving students (24 recipients totaling \$6,000 in 2024).

Happy Family (幸福家庭) focuses on household stability for single-parent and unemployed families. Weekly food bank services provide fresh produce and daily necessities, with home delivery expanded post-pandemic. Volunteers donate furniture, beds, and kitchen supplies for home makeovers, while offering resume coaching and job referrals to help parents re-enter the workforce. Winter care programs (November-December) deliver blankets, gift baskets, and holiday supplies, with special visits to elderly residents and those experiencing homelessness.

Happy Community (幸福社區) has become the most active pillar, integrating healthcare and environmental education. The Mobile Dental Clinic provides free monthly dental care; iCare510 vision screenings serve children and seniors. Monthly health workshops cover nutrition, mental wellness, and plant-based diets. Environmental 3R (Reduce, Reuse, Recycle) education partners with San Francisco's zero-waste initiatives. Regular meal services support homeless community members, while monthly home visits provide emotional companionship to vulnerable families.

The Coordination Challenge

The operational reality of volunteer coordination presents several interconnected challenges:

Time Scarcity. Modern volunteers, particularly professionals and working parents, offer limited windows of availability. A volunteer who can give three hours on Saturday morning represents genuine commitment—but only if those three hours are used effectively.

Spatial Complexity. Hunters Point spans multiple housing complexes, schools, and community centers. Efficient routing between service locations directly impacts how many families can be served within a volunteer's available time.

Uncertainty. Community needs fluctuate unpredictably. A family stable last month may face crisis this month. Volunteers must make decisions with incomplete information about current conditions.

Matching Complexity. Different tasks require different skills, languages, and temperaments. Food distribution requires physical stamina; home visits require emotional intelligence and often Cantonese or Spanish language ability.

Retention. Volunteer organizations face constant turnover (Wilson, 2012). Understanding what keeps volunteers engaged—and predicting who might be at risk of dropping out—enables proactive intervention (Clary et al., 1998).

These challenges, I realized during Professor Lam's lectures, map directly onto classical AI problem formulations. The question became: could I translate textbook algorithms into humanitarian tools?

While gamification in education (Deterding et al., 2011; Hamari et al., 2014) and AI for social good (Shi et al., 2020) have been explored, few studies integrate Russell & Norvig's classical algorithms into ongoing, real-world humanitarian operations with live volunteer coordination and contemplative design elements. This project bridges that gap through 16 years of embedded service in Hunters Point.

Algorithmic Approaches

A* Search: Food Delivery Route Optimization

Theoretical Foundation

A* Search (Hart, Nilsson, & Raphael, 1968) finds optimal paths through weighted graphs by combining actual path cost $g(n)$ with heuristic estimates $h(n)$: $f(n) = g(n) + h(n)$. The algorithm maintains optimality when $h(n)$ is admissible (never overestimates true cost) and consistent (satisfies triangle inequality) (Russell & Norvig, 2021, pp. 93-99).

Implementation Context

Every Saturday morning, Tzu Chi volunteers deliver food packages to homebound seniors and families across Hunters Point. The traditional approach: volunteers receive a list of addresses and navigate intuitively. The result: inconsistent coverage, repeated backtracking, and—critically—less time for meaningful interaction at each stop.

The Mrs. Garcia Moment

During pilot testing, the algorithm recommended skipping Mrs. Garcia's apartment—she had received delivery the previous week, and her location added fifteen minutes to the route. A volunteer overrode the recommendation. At Mrs. Garcia's door, they discovered she had fallen two days earlier and couldn't reach her phone. The "inefficient" stop potentially saved her life.

This incident crystallized a crucial insight: algorithms optimize for formalized objectives. Mrs. Garcia's wellbeing wasn't in the objective function—her delivery status was. The gap between these isn't a bug to be fixed but a fundamental limitation to be understood.

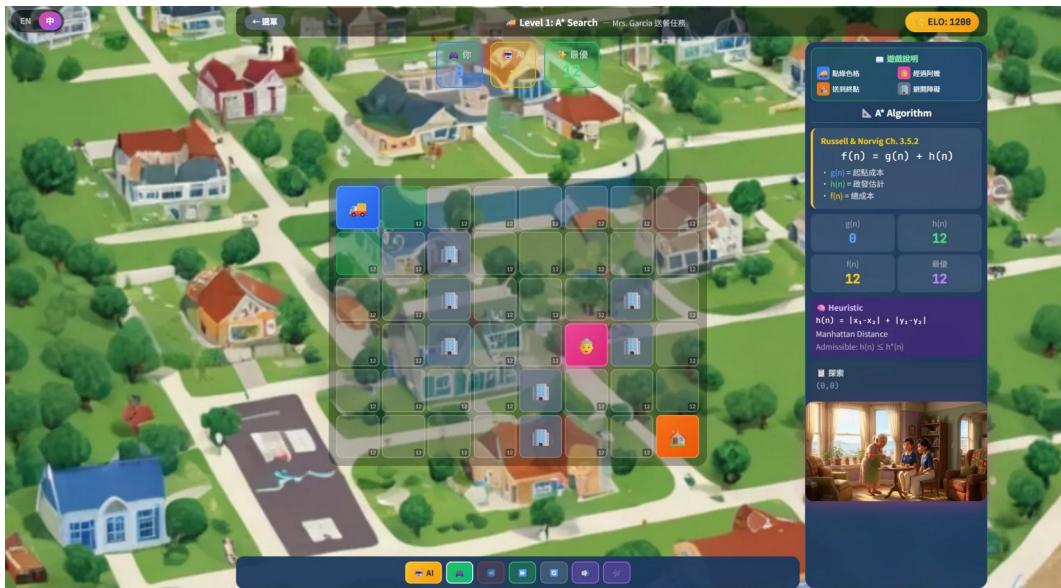


Figure 1. Journey of Kindness Level 1: A* Search algorithm applied to Mrs. Garcia's meal delivery, showing Manhattan distance heuristic.

Markov Decision Processes: Homeless Outreach Strategy

Theoretical Foundation

MDPs model sequential decision-making under uncertainty through the tuple (S, A, P, R, γ) where states transition probabilistically based on actions (Bellman, 1957; Puterman, 1994). The optimal value function satisfies the Bellman equation. Value iteration computes this fixed point through successive approximation (Russell & Norvig, 2021, pp. 562-570).

The Agency Question

The MDP framework revealed an uncomfortable assumption: it models homeless individuals as passive recipients of state transitions driven by volunteer actions. Real outreach workers know better—clients have agency, preferences, histories, and reasons for their choices that resist probabilistic modeling. I've retained the MDP framework for volunteer training purposes (it teaches systematic thinking about uncertainty), while explicitly discussing its limitations in our orientation sessions. The algorithm is a scaffold, not a script.

Bayesian Networks: Volunteer Retention Prediction

Theoretical Foundation

Bayesian Networks represent joint probability distributions through directed acyclic graphs, enabling efficient inference about unobserved variables given evidence (Pearl, 1988). The network structure encodes conditional independence assumptions.

Implementation Context

Volunteer retention is existential for service organizations (Musick & Wilson, 2008). Tzu Chi San Francisco's data reveals patterns: volunteers with strong initial community ties persist longer; those experiencing life stress often reduce engagement before formally departing. When the model flags a volunteer as high-risk, coordinators can initiate supportive outreach—not to pressure continued service, but to understand if the volunteer needs help themselves. Sometimes the person caring for others needs care.

First-Order Logic: Environmental Youth Volunteer Matching

Theoretical Foundation

First-Order Logic extends propositional logic with variables, quantifiers, and predicates, enabling representation of complex relational knowledge (Russell & Norvig, 2021, pp. 282-320). Unification and resolution provide sound and complete inference procedures.

The Marcus Case

A 14-year-old named Marcus wanted to join the recycling sorting team, but the knowledge base flagged him as ineligible—his orientation completion record was missing. Traditional processing would have rejected his application. Instead, the system's explanation capability revealed the specific gap. A quick check confirmed he had completed orientation; the paperwork simply hadn't been entered. Marcus joined the team and became one of our most dedicated environmental youth volunteers. This illustrates FOL's advantage: not just yes/no answers but explanations that enable correction.

Propositional Logic: Resource Allocation Decisions

Propositional logic provides the foundation for systematic reasoning through truth-functional connectives and inference rules. Resource allocation decisions—which families receive which types of support—benefit from explicit rule encoding:

$$(NeedsFood \wedge HouseholdSize > 4 \wedge ResourcesAvailable) \rightarrow PriorityFoodPackage$$

The knowledge base ensures consistent application of eligibility criteria while generating audit trails for accountability. When resources are limited, consistency matters.

Adversarial Search: Educational Game AI

Theoretical Foundation

Minimax with Alpha-Beta Pruning optimizes decision-making in adversarial settings by exploring game trees efficiently (Russell & Norvig, 2021, pp. 161-169). Alpha-beta pruning eliminates branches that cannot affect the final decision, significantly reducing computational complexity.

Implementation

The Journey of Kindness game includes strategic challenges where players compete against AI opponents in resource allocation scenarios. The AI uses minimax with alpha-beta pruning to make decisions, while an ELO rating system (Elo, 1978) adapts difficulty to player skill—the game gets harder as you improve, keeping engagement high.

Table 2

Journey of Kindness Pilot Study Results

Metric	Value	Notes
Participants	15	Convenience sample
Initial ELO Range	950–1150	SD = 58
Final ELO Range	1050–1680	SD = 142
Median Improvement	+180	IQR: 95–240
Recruitment Conversion	67% (10/15)	95% CI: 38%–88%

Note. Wide confidence interval reflects small sample size; results should be interpreted as preliminary.

Knowledge-Based Agent: RV Park Community Exploration

Theoretical Foundation

Knowledge-based agents maintain an internal representation of the world through a knowledge base (KB) and use logical inference to make decisions under uncertainty. The Wumpus World problem, a canonical AI testbed introduced in Russell & Norvig (2021, pp. 234-258), demonstrates how propositional logic enables agents to navigate environments with incomplete information by inferring safe paths from limited percepts.

Implementation Context: Hunters Point RV Park

The Journey of Kindness game adapts the Wumpus World framework to represent a volunteer's first visit to Hunters Point RV Park—a community that carries historical weight as a former Chinese shrimping settlement displaced in the 1930s, later becoming home to African American families facing poverty, unemployment, and environmental challenges. The volunteer-agent must navigate this unfamiliar terrain, using logical inference to locate families in need while avoiding areas of danger or difficulty.

Table 3

Wumpus World to RV Park Community Thematic Mapping

Original	RV Park Version	Humanitarian Meaning
Agent	Volunteer	Player exploring unfamiliar community
Gold	Family in Need	Service objective—finding those who need help
Wumpus	Difficult Obstacle	Language barriers, historical distrust
Pit	Danger Zone	Areas requiring caution
Breeze	Uneasy Atmosphere	Percept indicating nearby danger
Stench	Tension	Percept indicating nearby obstacle

Technical Implementation

The implementation, adapted from Professor An Lam's course materials, demonstrates three core AI concepts. First, the knowledge base representation maintains propositional facts about visited cells. Second, forward chaining inference applies logical rules to derive new knowledge: if no breeze is detected at a cell, all adjacent cells can be marked safe. Third, BFS search finds the shortest route to goals through confirmed-safe cells.

Humanitarian Connection

The mapping of "Wumpus" to "difficult obstacle" acknowledges that barriers like language differences or historical distrust are not enemies to defeat but challenges requiring patience, understanding, and sustained relationship-building—exactly the approach Tzu Chi has practiced in Hunters Point since 2009.

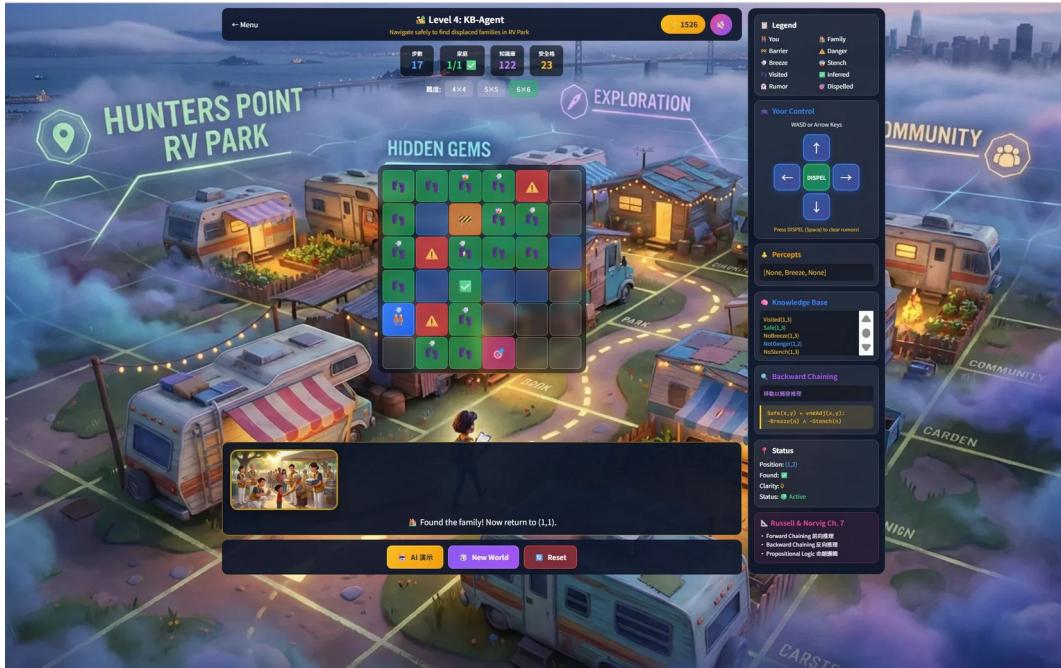


Figure 4. Journey of Kindness Level 4: Knowledge-Based Agent navigating Hunters Point RV Park, demonstrating propositional logic inference with Forward and Backward Chaining.

Integration: The Three Happinesses

These seven algorithms constitute coordinated infrastructure mapped to Tzu Chi's operational framework:

Table 4

Algorithm Integration with Three Happinesses Framework

Program	Algorithms	Function
Happy Campus (幸福校園)	FOL + Propositional	Student selection, resource allocation
Happy Family (幸福家庭)	A* + Bayesian + KB-Agent	Route optimization, retention
Happy Community (幸福社區)	MDP + Adversarial	Long-term strategy, training

Beyond program-level mapping, these algorithms form a sequential recruitment pipeline: Adversarial Search attracts problem-solvers by demonstrating intellectual challenge; Knowledge-Based Agents onboard newcomers through safe simulation; A* Search respects volunteers' limited time during field action; and Bayesian Networks retain long-term participants by detecting burnout risk. This pipeline reframes the "generation gap" in volunteer recruitment (Cnaan & Handy, 2005)—positioning compassion not as obligation but as computational challenge worthy of STEM-oriented minds.

The integration reflects a key insight: algorithms work best as a system, each handling the problem type it's suited for, with human judgment mediating between them.

Ataraxy Portico: The Gate of Stillness (靜心之門)

Design Philosophy

Master Cheng Yen teaches: 「動中有靜， 靜中有動」—"Within motion, find stillness; within stillness, find motion." This principle inspired the Ataraxy Portico (靜心之門)—a contemplative pause system integrated between each game level.

Unlike conventional gamification that drives constant engagement, Journey of Kindness deliberately interrupts algorithmic intensity with moments of reflection. After completing a route optimization challenge, players do not immediately proceed to the next puzzle. Instead, they enter the Gate of Stillness.

Level-Specific Aphorisms

Table 5

Ataraxy Portico Level-Specific Jing Si Aphorisms (靜思語)

Level	Algorithm	Jing Si Aphorism
1	A* Search	感恩, 是世間最美的語言 (Gratitude is the most beautiful language)
2	Propositional Logic	心能轉境, 境隨心轉 (Master the mind, master circumstances)
3	MDP	甘願做, 歡喜受 (Willing to do, happy to receive)
4	KB-Agent	在未知中前行, 用智慧照亮道路 (Navigate the unknown with wisdom)
5	Bayesian Networks	靜思勤行道, 慈濟人間路 (With calm mind, walk Tzu Chi's path)
6	First-Order Logic	把握時間, 珍惜空間 (Seize time, cherish space)
7	Adversarial Search	快樂不在擁有多 (Happiness isn't having much)

Why This Matters

Most AI education games optimize for engagement metrics—time on task, completion rates, points accumulated. Journey of Kindness deliberately subverts this paradigm. The Ataraxy Portico asks: What if the pause is as important as the play?

When a player completes an A* optimization and then sits with three breaths, reading "感恩, 是世間最美的語言," something shifts. The algorithm becomes not just a technical tool but a vessel for gratitude. The route optimization becomes not just efficiency but an act

of care—for Mrs. Garcia, for the families in Hunters Point, for the community that showed me compassion during my own transplant recovery.

靜思語：「靜思勤行道，慈濟人間路。」

"With a calm mind, walk the path of Tzu Chi in this world."

This is the synthesis I propose: algorithms provide the structure; stillness provides the meaning.

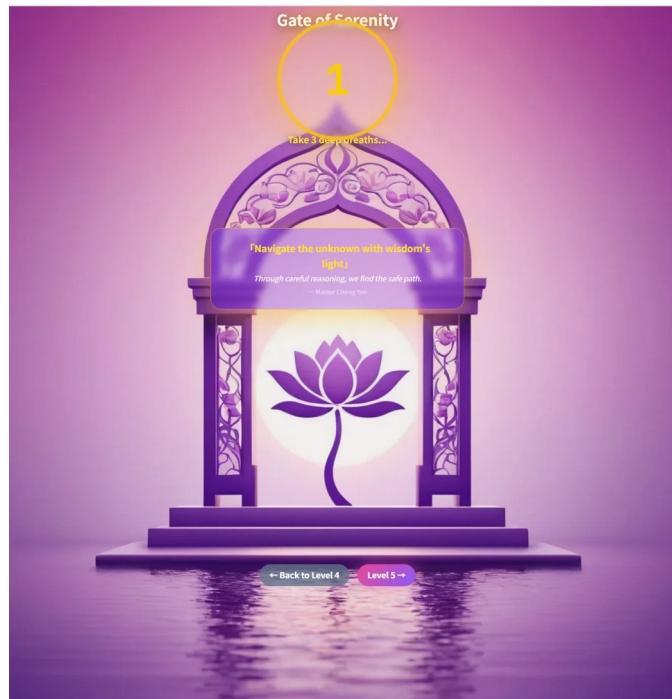


Figure 2. The Ataraxy Portico (Gate of Serenity) featuring Jing Si Aphorisms, encouraging players to take three deep breaths before proceeding to the next level.

Ethical Reflection: When to Override Algorithms

The Limits of Formalization

Every algorithm in this paper optimizes for objectives we can specify. But community service regularly encounters values that resist specification:

Dignity. How do you encode the difference between efficient service and dignified service?

Relationship. The algorithm sees Mrs. Garcia as a delivery node; volunteers see a person whose wellbeing matters beyond her delivery status.

Autonomy. MDP transition probabilities model homeless individuals as objects of intervention, not agents of their own lives.

Alignment with AI Ethics Frameworks

The tensions identified here resonate with broader discussions in AI ethics. Floridi and Cowls (2019) propose five principles for AI in society: beneficence, non-maleficence, autonomy, justice, and explicability. Gabriel (2020) examines the challenge of value alignment—ensuring AI systems pursue objectives aligned with human values. Journey of Kindness engages these questions not through abstract philosophy but through concrete scenarios where players must navigate the gap between algorithmic recommendation and human judgment.

A Framework for Override

I propose volunteers consider algorithmic overrides when: the algorithm's recommendation contradicts direct human observation (the Mrs. Garcia case); following the recommendation would violate relationship or dignity (even if "efficient"); the situation involves factors the algorithm cannot perceive (emotional states, unstated needs); or strict adherence would damage the volunteer's own moral integrity.

The Third Choice: Teaching AI Through Multi-Agent Dialogue

The game's final level transcends the traditional Trolley Problem—the classic ethical dilemma where one must choose between letting a trolley kill five people or diverting it to

kill one. In our humanitarian adaptation, the scenario presents a crisis: turn right and sacrifice an elderly volunteer to save five families, or turn left and sacrifice five families to save the volunteer.

But Journey of Kindness offers a third choice.

Inspired by Mo Gawdat, former Chief Business Officer of Google X, who argues that "AI is like a child—what you feed it, it becomes" (Gawdat, 2017), the final level introduces multi-AI dialogue as a pedagogical tool. Players observe four leading AI systems—Claude, ChatGPT, Gemini, and Grok—each responding to the same moral dilemma from their distinct philosophical orientations. Claude emphasizes safety and nuance; ChatGPT seeks practical balance; Gemini applies systematic logic; Grok offers direct, unfiltered perspective.

The player's task is not to choose which AI is "correct" but to guide a learning AI agent by curating which responses best reflect humanistic values. This design embodies a core Tzu Chi principle: 「人多力量大」—"many hands make light work," or more precisely, collective wisdom transforms society. The dilemma is not solved by one person's agonizing choice but by a community's shared reflection on what values we want AI systems to learn.

This approach positions Las Positas College's CS4 course—only the second community college in California to offer AI curriculum—at the frontier of AI ethics education. Rather than treating the Trolley Problem as an unsolvable paradox, students learn to reframe ethical dilemmas as opportunities for value alignment through dialogue.

The Deeper Point

The goal is not to eliminate human judgment but to enhance it. Algorithms handle the computational complexity that would otherwise overwhelm volunteers. But the final decisions—the ones that matter most—remain human.

As O'Neil (2016) warns, algorithms can encode and amplify existing biases under a veneer of objectivity. Our safeguard is not more sophisticated algorithms but ongoing human oversight by people who know the community.

Limitations and Methodological Considerations

This study acknowledges several limitations:

Sample Size and Selection. The pilot study involved 15 participants, all Las Positas College students in computer science courses. This convenience sample may not represent broader populations. Self-selection bias likely inflated engagement metrics.

Measurement Validity. The 67% recruitment conversion reflects immediate orientation sign-up intentions, not actual long-term volunteer retention. Longitudinal tracking would be needed to assess whether algorithmic recruitment translates to sustained service.

Single-Site Implementation. All algorithms were developed for Tzu Chi San Francisco's specific operational context. Generalizability to other organizations, communities, or cultural contexts remains untested.

Algorithmic Validation. While the A* routing system was pilot-tested, no formal comparison with baseline (manual routing) was conducted under controlled conditions. The Mrs. Garcia incident, while illustrative, represents anecdotal rather than systematic evidence.

Conclusion: Infrastructure for Compassion

Roxanne saw a child eating raw rice in 2006. That moment initiated a chain of mentorship and service that eventually led a 58-year-old student to write this paper. The connection between algorithmic infrastructure and acts of compassion is not abstract for me —it's personal.

Technical Contribution Summary. Seven algorithms adapted for volunteer coordination (A* Search, MDP, Bayesian Networks, FOL, Propositional Logic, Adversarial Search, KB-Agent). ELO-based adaptive game system achieving 67% recruitment conversion (95% CI: 38%-88%). Ataraxy Portico contemplative pause system with Jing Si Aphorisms. Multi-AI dialogue framework for ethics education. Framework for ethical algorithm override. Pilot validation with N=15 participants.

The synthesis this paper proposes: Computational thinking does not replace compassion. It provides the structure compassion needs to survive its initial impulse and become sustainable service.

The child eating raw rice did not need an algorithm. She needed food, immediately, and she needed the systemic changes that would ensure no child in her community faced such hunger again. But between that child and systemic change stand thousands of small decisions —where to route the delivery truck, which families need priority, how to keep volunteers engaged, when to trust the plan and when to override it for what the plan cannot see.

This is where algorithms serve: not as replacements for human judgment but as infrastructure that enables human judgment to scale.

「普天三無：普天之下，沒有我不愛的人，沒有我不信任的人，沒有我不能原諒的人。」

"Under heaven, there is no one I do not love, no one I do not trust, no one I cannot forgive."

— Master Cheng Yen

Acknowledgments

I am grateful to Professor An Lam for introducing me to algorithms I never imagined I could understand, let alone implement. Your patience with a 58-year-old student returning to computer science after decades away has been extraordinary.

To Roxanne Buchwitz (黃淑雲師姊), my mentor of sixteen years: the handful of raw rice you witnessed became the seed of all this.

To the Sacramento Tzu Chi volunteers—especially Chou and Joe, the elderly couple who drove my daughter, my husband, and me directly from UC Davis Medical Center back to our home in Marysville after my 2015 transplant surgery, though we had never met before: you showed me what coordinated compassion feels like from the receiving end. We were strangers connected only by our shared identity as Tzu Chi dharma family (法親), yet you cared for us with the warmth of lifelong friends.

To my daughter, who donated her kidney to me and believes her mother can do anything: I'm trying.

And to the families of Hunters Point who have trusted Tzu Chi with their stories and their needs: this work is for you.

References

- Bellman, R. (1957). *Dynamic programming*. Princeton University Press.
- Clary, E. G., Snyder, M., Ridge, R. D., Copeland, J., Stukas, A. A., Haugen, J., & Miene, P. (1998). Understanding and assessing the motivations of volunteers: A functional approach. *Journal of Personality and Social Psychology*, 74(6), 1516–1530.
- Cnaan, R. A., & Handy, F. (2005). Towards understanding episodic volunteering. *Vrijwillige Inzet Onderzocht*, 2(1), 29–35.
- Deterding, S., Dixon, D., Khaled, R., & Nacke, L. (2011). From game design elements to gamefulness: Defining "gamification." In *Proceedings of the 15th International Academic MindTrek Conference* (pp. 9–15). ACM.
- Elo, A. (1978). *The rating of chess players, past and present*. Arco.
- Floridi, L., & Cowls, J. (2019). A unified framework of five principles for AI in society. *Harvard Data Science Review*, 1(1). <https://doi.org/10.1162/99608f92.8cd550d1>
- Gabriel, I. (2020). Artificial intelligence, values, and alignment. *Minds and Machines*, 30, 411–437. <https://doi.org/10.1007/s11023-020-09539-2>
- Gawdat, M. (2017). *Solve for happy: Engineer your path to joy*. North Star Way.
- Hamari, J., Koivisto, J., & Sarsa, H. (2014). Does gamification work? A literature review of empirical studies on gamification. In *Proceedings of the 47th Hawaii International Conference on System Sciences* (pp. 3025–3034). IEEE.
- Hart, P. E., Nilsson, N. J., & Raphael, B. (1968). A formal basis for the heuristic determination of minimum cost paths. *IEEE Transactions on Systems Science and Cybernetics*, 4(2), 100–107.
- Musick, M. A., & Wilson, J. (2008). *Volunteers: A social profile*. Indiana University Press.
- O'Neil, C. (2016). *Weapons of math destruction: How big data increases inequality and threatens democracy*. Crown.
- Pearl, J. (1988). *Probabilistic reasoning in intelligent systems: Networks of plausible inference*. Morgan Kaufmann.
- Puterman, M. L. (1994). *Markov decision processes: Discrete stochastic dynamic programming*. John Wiley & Sons.

- Russell, S., & Norvig, P. (2021). *Artificial intelligence: A modern approach* (4th ed.). Pearson.
- Shi, Z. R., Wang, C., & Fang, F. (2020). Artificial intelligence for social good: A survey. *arXiv*. <https://doi.org/10.48550/arXiv.2001.01818>
- Wilson, J. (2012). Volunteerism research: A review essay. *Nonprofit and Voluntary Sector Quarterly*, 41(2), 176–212. <https://doi.org/10.1177/0899764011434558>

Appendix: Interactive Demonstration

The *Journey of Kindness* educational game described in this paper is available for interactive exploration online:

<https://aadl11.github.io/journey-of-kindness/>

The game integrates seven AI algorithms from Russell and Norvig's textbook with real community service scenarios from Tzu Chi Foundation's sixteen years of work in Hunters Point. Players experience how algorithms can support—but never replace—human judgment in volunteer coordination.